

Operating System

Assignment 3

Eran Aflalo 209343722

Omer Luxembourg 205500390

1 Signal Framework

We implemented a signal handler that can receive 2 numbers – pid of the sending process and a number that is sent by the sending process.

We defined a typedef which is a declaration of a signal handler function:

```
types.h  proc.h  proc.c  exec.c  user.h  usys.S
1  typedef unsigned int uint;
2  typedef unsigned short ushort;
3  typedef unsigned char uchar;
4  typedef uint pde_t;
5  //----- 06.01.21 -----
6  typedef void (*sig_handler)(int pid, int value);
7  //-----
```

In addition we defined the system call function **sigset** and **sigsend** and **sigret** (for the user space also...):

```
trapasm.S  types.h  proc.h  proc.c  exec.c  user.h  usys.S  syscall.h  syscall.c
27  //----- 06.01.21 -----
28  //set the signal handler to be called when signals are sent
29  sig_handler sigset(sig_handler);
30  // send a signal with the given value to a process with pid dest_pid
31  int sigsend(int dest_pid, int value);
32  //----- 09.01.21 -----
33  void sigret(void);
34  //-----
```

Sigret will be mentioned in the 4th part.

The three of them been added as in task 1, to the system call chain. (**usys.S**, **syscall.h**, **syscall.c**, etc)

System call route:

```
trapasm.S  types.h  proc.h  usys.S  syscall.c  defs.h  syscall.h  proc.c
103  //PAGEBREAK: 16
104  // proc.c
105
106  int    cpuid(void);
107  void   exit(void);
108  int    fork(void);
109  int    growproc(int);
110  int    kill(int);
111  struct cpu* mycpu(void);
112  struct proc* myproc();
113  void   pinit(void);
114  void   procdump(void);
115  void   scheduler(void) __attribute__((noreturn));
116  void   sched(void);
117  void   setproc(struct proc*);
118  void   sleep(void*, struct spinlock*);
119  void   userinit(void);
120  int    wait(void);
121  void   wakeup(void*);
122  void   yield(void);
123  //----- 06.01.21 -----
124  int    sigsend(int,int);
125  //----- 09.01.21 -----
126  void   sigret(void);
127  //-----
```

```
trapasm.S  types.h  syscall.h  proc.h
18  #define SYS_mknod 17
19  #define SYS_unlink 18
20  #define SYS_link 19
21  #define SYS_mkdir 20
22  #define SYS_close 21
23
24  //----- 06.01.21 -----
25  #define SYS_sigset 22
26  #define SYS_sigsend 23
27  //----- 09.01.21 -----
28  #define SYS_sigret 24
29  //-----
```

```
trapasm.S  types.h  proc.h  usys.S  proc.c
28  SYSCALL(dup)
29  SYSCALL(getpid)
30  SYSCALL(sbrk)
31  SYSCALL(sleep)
32  SYSCALL(uptime)
33
34  #----- 06.01.21 -----
35  SYSCALL(sigset)
36  SYSCALL(sigsend)
37  #----- 09.01.21 -----
38  SYSCALL(sigret)
```

```
trapasm.S  types.h  proc.h  usys.S  syscall.c
103  extern int sys_wait(void);
104  extern int sys_write(void);
105  extern int sys_uptime(void);
106  //----- 06.01.21 -----
107  extern int sys_sigset(void);
108  extern int sys_sigsend(void);
109  //----- 09.01.21 -----
110  extern int sys_sigret(void);
111  //-----
```

```
trapasm.S  types.h  proc.h  usys.S  syscall.c
131  [SYS_mknod] sys_mknod,
132  [SYS_unlink] sys_unlink,
133  [SYS_link] sys_link,
134  [SYS_mkdir] sys_mkdir,
135  [SYS_close] sys_close,
136
137  //----- 06.01.21 -----
138  [SYS_sigset] sys_sigset,
139  [SYS_sigsend] sys_sigsend,
140  //----- 09.01.21 -----
141  [SYS_sigret] sys_sigret,
142  //-----
143  };
144
```

The rest of the implementation is in the next sections.

2 Storing and Changing the Signal Handler

To store the signal handler, we will add a new field to struct proc (in **proch.h**). This field will hold a pointer to the current handler (or -1 if no handler is set). Both **fork** and **exec** system calls modify the signal handler.

fork –

```
187 // Caller must set state of returned proc to RUNNABLE.
188 int
189 fork(void)
190 {
191     int i, pid;
192     struct proc *np;
193     struct proc *curproc = myproc();
194
195     // Allocate process.
196     if((np = allocproc()) == 0){
197         return -1;
198     }
199
200     // Copy process state from proc.
201     if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
202         kfree(np->kstack);
203         np->kstack = 0;
204         np->state = UNUSED;
205         return -1;
206     }
207     np->sz = curproc->sz;
208     np->parent = curproc;
209     *np->tf = *curproc->tf;
210
211     // Clear %eax so that fork returns 0 in the child.
212     np->tf->eax = 0;
213
214     for(i = 0; i < NOFILE; i++)
215         if(curproc->ofile[i])
216             np->ofile[i] = filedup(curproc->ofile[i]);
217     np->cwd = idup(curproc->cwd);
218     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
219
220     pid = np->pid;
221
222     acquire(&ptable.lock);
223     //----- 06.01.21 -----
224     np->sigHandler = curproc->sigHandler;
225     //-----
226     np->state = RUNNABLE;
227     release(&ptable.lock);
228
229
230
231     return pid;
}
```

Takes parent handler

exec –

exec body function

default handler (-1)

```
74     goto bad;
75     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
76     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
77         goto bad;
78     ustack[3+argc] = sp;
79 }
80 ustack[3+argc] = 0;
81
82 ustack[0] = 0xffffffff; // fake return PC
83 ustack[1] = argc;
84 ustack[2] = sp - (argc+1)*4; // argv pointer
85
86 sp -= (3+argc+1) * 4;
87 if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
88     goto bad;
89
90 // Save program name for debugging.
91 for(last=s=path; *s; s++)
92     if(*s == '/')
93         last = s+1;
94 safestrcpy(curproc->name, last, sizeof(curproc->name));
95
96 // Commit to the user image.
97 oldpgdir = curproc->pgdir;
98 curproc->pgdir = pgdir;
99 curproc->sz = sz;
100 curproc->tf->eip = elf.entry; // main
101 curproc->tf->esp = sp;
102 //----- 06.01.21 -----
103 curproc->sigHandler = (sig_handler)(-1);
104 //-----
105 switchvm(curproc);
106 freevm(oldpgdir);
107 return 0;
108
109 bad:
110 if(pgdir)
111     freevm(pgdir);
112 if(ip){
113     iunlockput(ip);
114     end_op0;
115 }
116 return -1;
117 }
118
```

Sigset system call (a.k.a. **sys_sigset**) is defined in **sysproc.c** to replace the destination process' signal handler with a new one, and return the old signal handler.

```
93 //----- 06.01.21 -----
94 //system call - replace the old signal handler to new signal handler and return the old one.
95 int
96 sys_sigset(void)
97 {
98     sig_handler new_sigHandler;
99     if(argptr(0, (void*)&new_sigHandler, sizeof(sig_handler)) < 0)
100         return -1;
101     sig_handler old_sigHandler = myproc()->sigHandler;
102     myproc()->sigHandler = new_sigHandler;
103     cprintf("sysproc.c : sys_sigset nom. for process %d\n", myproc()->pid);
104
105     return (int)(old_sigHandler);
106 }
107
```

3 Sending a Signal to a Process

The new **sigsend** system call sends a signal to a destination process. When a signal is sent to a process it is not handled instantly since the destination process may be already running or even blocked. This means that each process must store all the signals which were sent to it but still not handled in a data structure that we will refer to as the pending signals stack, Since multiple processes can send signals to the same recipient then he must save the signals in structure.

The **sigsend** system call will add a record to the recipient pending signals stack. It will return 0 on success and -1 on failure (if pending signals stack is full). Systemcall function added to **sysproc.c**:

```
107
108 //system call - send signal "value" from current processs to dest_pid process.
109 int
110 sys_sendsend(void)
111 {
112     int dest_pid;
113     int value;
114     if(argint(0, &dest_pid) < 0)
115         return -1;
116     if(argint(1, &value) < 0)
117         return -1;
118     cprintf("(sysproc.c) sys_sendsend: from process %d - to pid %d with value %d.\n", myproc()->pid, dest_pid, value);
119     return sigsend(dest_pid, value);
120 }
```

For each process, we created the following –

```
1 //----- 06.01.21 -----
2
3 // defines an element of the concurrent struct
4 struct cstackframe{
5     int sender_pid;
6     int recipient_pid;
7     int value;
8     int used;
9     struct cstackframe *next;
10 };
11
12 //defines a concurrent stack
13 struct cstack{
14     struct cstackframe frames[10];
15     struct cstackframe *head;
16 };
17
18 // adds a new frame tot the cstack which is initialized with values send_pid,
19 // recipient_pid and value,
20 // then returns 1 on success and 0 if the stack is full
21 int push (struct cstack *cstack, int sender_pid, int recipient_pid, int value);
22
23 // removes and returns an element from the head of given cstack. if the stack is empty, then return 0
24 struct cstackframe *pop(struct cstack *cstack);
25
26 //-----end-----
27
```

```

60 };
61
62 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
63
64 // Per-process state
65 struct proc {
66     uint sz;           // Size of process memory (bytes)
67     pde_t* pgdir;      // Page table
68     char* kstack;      // Bottom of kernel stack for this process
69     enum procstate state; // Process state
70     int pid;           // Process ID
71     struct proc* parent; // Parent process
72     struct trapframe* tf; // Trap frame for current syscall
73     struct context* context; // switch() here to run process
74     void* chan;        // If non-zero, sleeping on chan
75     int killed;        // If non-zero, have been killed
76     struct file*ofile[NOFILE]; // Open files
77     struct inode* cwd; // Current directory
78     char name[16];     // Process name (debugging)
79
80     //----- 06.01.21 -----
81     sig_handler sigHandler;
82     struct cstack pending_signals;
83     //-----
84
85 };
86

```

We initialized in **proc.c** – **allocproc()** function the cstack pending_signals frames and head:

```

114
115 //----- 07.01.21 -----
116 //reset the used of the pending_signals and the head to &frames[0]
117 int i;
118 for(i=0;i<10;i++){
119     (p->pending_signals).frames[i].used=0;
120     (p->pending_signals).head=&((p->pending_signals).frames[0]);
121 }
122 //-----
123 return p;
124 }

```

The **sigsend**, **push** and **pop** functions –

```

546 //----- 06.01.21 -----
547 int
548 sigsend(int dest_pid,int value)
549 {
550     struct proc* p;
551     acquire(&ptable.lock);
552     // finding the destination process
553     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
554         if(p->pid == dest_pid){
555             // try to push the signal to cstack of the destination process
556             int success = push(&(p->pending_signals),myproc()->pid,dest_pid,value);
557             release(&ptable.lock);
558             return (success-1); //returns 0 on success and -1 if failed (full cstack)
559         }
560     }
561     release(&ptable.lock);
562     return -1; //returns -1 because failed to find the process
563 }
564
565 // adds a new frame tot the cstack which is initialized with values send_pid,
566 // recipient_pid and value,
567 // then returns 1 on success and 0 if the stack is full
568 int push (struct cstack* cstack, int sender_pid, int recipient_pid, int value)
569 {
570     int i;
571     //make the new frame which will be the new head of the stack (head become second)
572     struct cstackframe new_csframe = {sender_pid,recipient_pid,value,1,cstack->head};
573     for(i=0; i<10;i++){
574         if((cstack->frames[i]).used==0){
575             cstack->frames[i] = new_csframe;
576             cstack->head = &(cstack->frames[i]);
577             return 1;
578         }
579     }
580     return 0; //if full
581     //if(cstack->head==null)
582 }
583
584
585

```

Sending the destination process a signal.

```

595 // removes and returns an element from the head of given cstack. if the stack is empty, then return 0
596 struct cstackframe* pop(struct cstack* cstack)
597 {
598     struct cstackframe* oldHead = cstack->head;
599     if(oldHead->used==0) //the stack is empty
600         return 0;
601     else{
602         //remove the head of the stack
603         (cstack->head)->used=0;
604         cstack->head = oldHead->next;
605         return oldHead;
606     }
607 }
608
609
610

```

4 Signal Handling

When a process is about to return from kernel space to user space (using the function **trapret** which can be found at **trapasm.S**) it must check its pending signals stack. If a pending signal exists and the process is not already handling a signal (i.e., we did not support handling multiple signals at once) then the process must handle the signal.

The signal handling can be either discarding the signal (if the signal handler is default we printed a message) or executing a signal handler when it returns to user space. To force the execution of the signal handler in user space we have to modified the user space stack and the instruction pointer of the process.

We changed the following to make the function call: pushed the arguments for the called function, pushed the return address on the stack, and jump to the body of the called function. In addition in the return address we pointed it to a code we injected to the stack – implicit sigret.

Also, we saved the CPU registers values and restored them when returning from the called function – signal handler.

This is the system call **sigret()**:

```
122 //----- 09.01.21 -----
123 //system call - sigret
124 void
125 sys_sigret(void)
126 {
127     sigret();
128     return;
129 }
130 //-----
```

In **proc.c** we created the **sigret** system call, which only restoring our backup of the **trapframe**.

```
612 //----- 10.01.21 -----
613 // a system call function to handle signal return - using the old trapframe and handling signal to 0
614 void
615 sigret(void)
616 {
617     struct proc *p = myproc();
618     memmove(p->tf, &p->tf_backup, sizeof(struct trapframe)); // restoring the old backup of the trapframe
619     p->handling_signal = 0; // turning off the handling signals flag - now other signals can be handled
620 }
```

In the **trapasm.S** we updated the stack to call the **signals_check** (which is the signals checker function). In addition we added an **implicit_sigret** to the file, which will be loaded to the user's stack, to execute after returning from **signals_check** function.

```
26 # Return falls through to trapret...
27 .globl trapret
28 trapret:
29 # ----- 06.01.21 -----
30 pushl %esp
31 call signals_check # check for pending signals, if there are, handle them if you available
32 addl $4, %esp
33 # -----
34 popal
35 popl %gs
36 popl %fs
37 popl %es
38 popl %ds
39 addl $0x8, %esp # trapno and errcode
40 iret
41
42
43 # ----- 10.01.21 -----
44 .global implicit_sigret_start
45 .global implicit_sigret_end
46 implicit_sigret_start: # inserting a call to the sigret sys_call
47 movl $SYS_sigret, %eax
48 int $T_SYSCALL # User code makes a system call with INT T_SYSCALL.
49 implicit_sigret_end:
50 # -----
```

The function **signals_check(tf)** receiving the old *trapframe* and backing it up for future resotre (as we mentioned in **sigret**). We handeled the signal which was given by popping the head frame from the *cstackframe* of the process – the handling will be the deafult (a simple printing function) if the *sig_handler* is -1. If the **sig_handler** is initialized we are executing it by pushing the arguments (**value** and **send_pid**) and the **sig_handler**. It will be executed after we are done with the function **signals_check**. The next PC after the execution of the handler will be the in the stack, where we entered the code for the **implicit_sigret**.

```

622 void signals_check(struct trapframe *tf){
623     struct proc *p = myproc();
624     if (p == 0) //this CPU has no proc defined
625         return;
626     if (p->handling_signal) // the process is handling a signal, return...
627         return;
628     struct cstackframe* head_frame = pop(&p->pending_signals);
629     if (head_frame == 0) // this popped frame is null
630         return;
631     if (p->sigHandler == (sig_handler)(-1)){ // this is the default signal handler
632         cprintf("(proc.c) signals_check: default sig_handler - sender_pid %d, recipient_pid %d, value %d\n", head_frame->sender_pid, head_frame->recipient_pid, head_frame->value);
633         return;
634     }
635     p->handling_signal = 1;
636     memmove(&p->tf_backup, p->tf, sizeof(struct trapframe)); //backup of the trapframe
637
638     uint length = (uint)&implicit_sigret_end - (uint)&implicit_sigret_start;
639     p->tf->esp -= length; // making room for variables
640     memmove((void*)p->tf->esp, implicit_sigret_start, length); // copying implicit call to sigret to the stack
641     *((int*)(p->tf->esp - 4)) = head_frame->value; // value -> to the stack
642     *((int*)(p->tf->esp - 8)) = head_frame->sender_pid; // sender_pid -> to the stack
643     *((int*)(p->tf->esp - 12)) = p->tf->esp; // implicit sigret sys_call return address
644     p->tf->esp -= 12;
645     p->tf->eip = (uint)p->sigHandler; // start the function sigHandler, right after the trapret
646     head_frame->used = 0; // realising the cstack
647     return;
648 }
649 //-----

```


5 Signal Handling

We created the following function in the user space:

```
7  ///passing command line arguments
8
9  void theBestSignalHandler(int pid, int value){
10     printf(1, "new sig_handler: I am the best signal handler my pid %d from recipient_pid %d, value %d\n", getpid(), pid, value);
11 }
12
13 int main(int argc, char **argv)
14 {
15     printf(1, "===== Default sig_handler =====\n");
16
17     int pid = fork();
18
19     if(pid == 0){
20         sleep(3);
21         printf(1, "Marco has woken up...\n");
22         exit();
23     } else if(pid < 0){
24         printf(1, "This is fork failed\n");
25     }
26     printf(1, "Anna (pid %d) is sending Marco (pid %d) a signal...\n", getpid(), pid);
27     sigsend(pid, 17);
28     wait();
29     printf(1, "Anna has done waiting for Marco...\n");
30
31     printf(1, "===== New sig_handler =====\n");
32     pid = fork();
33     if(pid == 0){
34         sigset((sig_handler)&theBestSignalHandler);
35         sleep(3);
36         printf(1, "Marco has woken up...\n");
37         exit();
38     } else if(pid < 0){
39         printf(1, "This is fork failed\n");
40     }
41     sleep(2);
42     printf(1, "Anna (pid %d) is sending Marco (pid %d) a signal...\n", getpid(), pid);
43     sigsend(pid, 17);
44     wait();
45     printf(1, "Anna has done waiting for Marco...\n");
46
47     printf(1, "===== END TESTING =====\n");
48     exit();
49 }
50
```

We tested the system with 2 signal handlers – the first one is the default signal handler (print only), and the second can be anything, we created *the best signal handler* to handler it.

```
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ sig_test
===== Default sig_handler =====
Anna (pid 3) is sending Marco (pid 4) a signal...
(sysproc.c) sys_sigsend: from process 3 - to pid 4 with value 17.
(proc.c) signals_check: default sig_handler - sender_pid 3, recipient_pid 4, value 17
Marco has woken up...
Anna has done waiting for Marco...
===== New sig_handler =====
(sysproc.c) sys_sigset: for process 5.
Anna (pid 3) is sending Marco (pid 5) a signal...
(sysproc.c) sys_sigsend: from process 3 - to pid 5 with value 17.
new sig_handler: I am the best signal handler my pid 5 from recipient_pid 3, value 17
Marco has woken up...
Anna has done waiting for Marco...
===== END TESTING =====
$
```

The results, are as we expected.